

自然语言处理

week-5

凌震华

2024年3月28日





□ Probabilistic CFGs

□ Structural Dependencies

□ Parser Evaluation



Why Statistical Parsing

- Solve problems of syntactic ambiguity
 - Parsing algorithms can identify it, but cannot resolve it
- Model human parsing which is known to be probabilistic
- Extend parsing to poorly understood languages for which data is available
- Efficient parsing for sub-languages restricted to specific domains



Probabilistic CFGs

- The probabilistic model
 - Assigning probabilities to parse trees
- Getting the probabilities for the model
- Parsing with probabilities
 - Slight modification to dynamic programming approach
 - Task is to find the max probability tree for an input



Probability Model

- Attach probabilities to grammar rules
- The expansions for a given non-terminal sum to 1

VP \rightarrow Verb .55

VP \rightarrow Verb NP .40

VP \rightarrow Verb NP NP .05

– Read this as $P(\text{Specific rule} \mid \text{LHS})$

- $P(\text{VP} \rightarrow \text{Verb} \mid \text{VP}) = .55$



Probability Model

- A derivation (tree) consists of the bag of grammar rules that are in the tree
- The probability of a tree is just the product of the probabilities of the rules in the derivation.

$$P(T, S) = \prod_{node \in T} P(rule(n))$$





Probability Model

- The probability of a word sequence (sentence) is the probability of its tree in the unambiguous case.
- It' s the sum of the probabilities of the trees in the ambiguous case.
- Since we can use the probability of the tree(s) as a proxy for the probability of the sentence...
 - PCFGs give us an alternative to N-Gram models as a kind of language model.



Example



Rule Probabilities

Rules			P	Rules			P
S	→	VP	.05	S	→	VP	.05
VP	→	Verb NP	.20	VP	→	Verb NP NP	.10
NP	→	Det Nominal	.20	NP	→	Det Nominal	.20
Nominal	→	Nominal Noun	.20	NP	→	Nominal	.15
Nominal	→	Noun	.75	Nominal	→	Noun	.75
Verb	→	book	.30	Nominal	→	Noun	.75
Det	→	the	.60	Verb	→	book	.30
Noun	→	dinner	.10	Det	→	the	.60
Noun	→	flights	.40	Noun	→	dinner	.10
				Noun	→	flights	.40

$2.2 * 10^{-6}$

$6.1 * 10^{-7}$



Getting the Probabilities

- From an annotated database (a treebank)
 - So for example, to get the probability for a particular VP rule just count all the times the rule is used and divide by the number of VPs overall.

$$P(\alpha \rightarrow \beta | \alpha) = \frac{\text{Count}(\alpha \rightarrow \beta)}{\sum_{\gamma} \text{Count}(\alpha \rightarrow \gamma)} = \frac{\text{Count}(\alpha \rightarrow \beta)}{\text{Count}(\alpha)}$$



Getting the Probabilities

- If we don't have a treebank, but we do have a grammar can we get reasonable probabilities?
- Yes. Use a prob parser to parse a large corpus and then get the counts as above.
 - In the unambiguous case we're fine
 - In ambiguous cases, weight the counts of the rules by the probabilities of the trees they occur in.
 - Where do those probabilities come from?
 - Make them up. And then re-estimate them.



Assumptions

- We're assuming that there is a **grammar** to be used to parse with.
- We're assuming the existence of a large robust **dictionary** with parts of speech
- We're assuming the ability to parse (i.e. **a parser**)
- Given all that... we can parse probabilistically



Typical Approach

- Use CKY as the backbone of the algorithm
- Assign probabilities to constituents as they are completed and placed in the table
- Use the max probability for each constituent going up



Clarifying last point...

- Say we're talking about a final part of a parse
 - $S \rightarrow_0 NP_i VP_j$

The probability of this S is...

$$P(S \rightarrow NP VP) * P(NP) * P(VP)$$

The green stuff is already known if we're using some kind of sensible DP approach.



Max

- I said the $P(NP)$ is known.
- What if there are multiple NPs for the span of text in question (0 to i)?
- Take the max



CKY

```
function CKY-PARSE(words, grammar) returns table

  for  $j \leftarrow$  from 1 to LENGTH(words) do
     $table[j-1, j] \leftarrow \{A \mid A \rightarrow words[j] \in grammar\}$ 
    for  $i \leftarrow$  from  $j-2$  downto 0 do
      for  $k \leftarrow i+1$  to  $j-1$  do
         $table[i, j] \leftarrow table[i, j] \cup$ 
           $\{A \mid A \rightarrow BC \in grammar,$ 
             $B \in table[i, k],$ 
             $C \in table[k, j]\}$ 
```



Prob CKY

function PROBABILISTIC-CKY(*words*,*grammar*) **returns** most probable parse
and its probability

```
for  $j \leftarrow$  from 1 to LENGTH(words) do
  for all  $\{ A \mid A \rightarrow words[j] \in grammar \}$ 
     $table[j-1, j, A] \leftarrow P(A \rightarrow words[j])$ 
  for  $i \leftarrow$  from  $j-2$  downto 0 do
    for  $k \leftarrow i+1$  to  $j-1$  do
      for all  $\{ A \mid A \rightarrow BC \in grammar,$ 
        and  $table[i, k, B] > 0$  and  $table[k, j, C] > 0 \}$ 
        if  $(table[i, j, A] < P(A \rightarrow BC) \times table[i, k, B] \times table[k, j, C])$  then
           $table[i, j, A] \leftarrow P(A \rightarrow BC) \times table[i, k, B] \times table[k, j, C]$ 
           $back[i, j, A] \leftarrow \{k, B, C\}$ 
return BUILD_TREE( $back[1, LENGTH(words), S]$ ),  $table[1, LENGTH(words), S]$ 
```



Problems with PCFGs

- The probability model we're using is just based on the rules in the derivation...
 - Doesn't take into account **where** in the derivation a rule is used (structural issue)
 - Doesn't use the words in any real way
 - Doesn't really work
 - Most probable parse isn't usually the right one (the one in the treebank test set).



□ Probabilistic CFGs

□ Structural Dependencies

□ Parser Evaluation



Structural dependencies

- Strength of CFG -- rules are “context-free” – becomes a weakness
- Cannot model context-sensitive probabilities
- See section 14.4.1
 - NP that is a subject is more likely to be a pronoun
 - NP → PRP (.91 in subject pos, .34 in object pos)
 - Context-free, it resolves to .25 probability



One solution

- Parent annotation 父结点标注 (section 14.5)
- Split non-terminals to add context
- Parent info added to every non-terminal
 - Subject NP becomes NP^S ; Object NP becomes NP^O
- Split pre-terminal POS nodes
 - Note that in parsing, POS categories are fixed and come from “outside” the grammar



Parent Annotation

- Increases size/complexity of grammar
- Requires more data than is typically available
- Split-and-merge algorithm now available for automatically adjusting the size
- Gives the best published result on penn-treebank.



Solution for using word info

- Add lexical dependencies to the scheme...
 - Integrate the preferences of particular words into the probabilities in the derivation
 - i.e. Condition the rule probabilities on the actual words

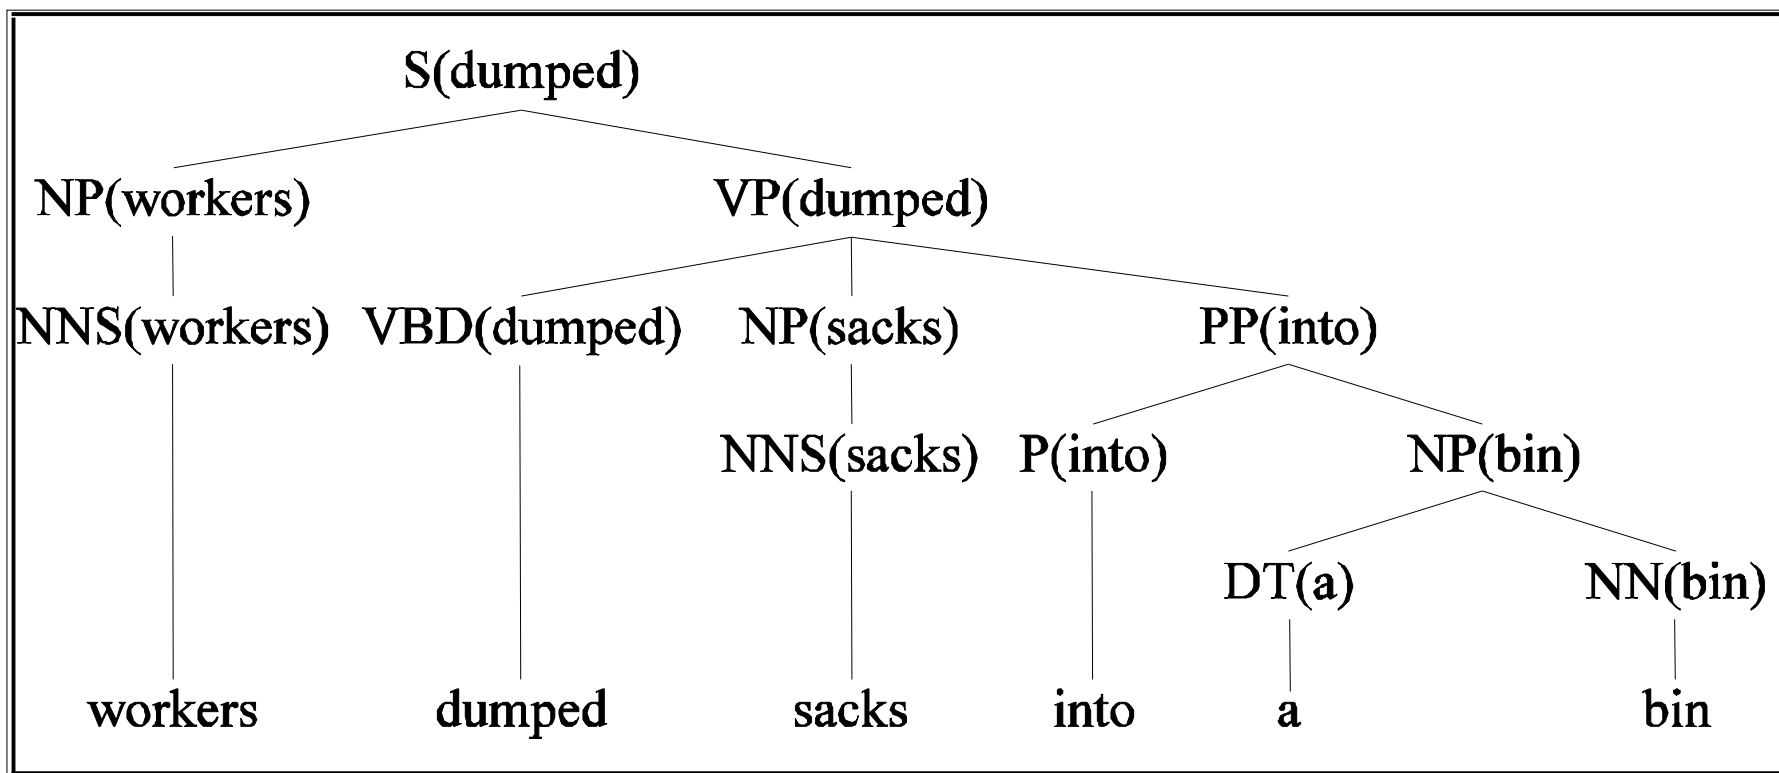


Heads

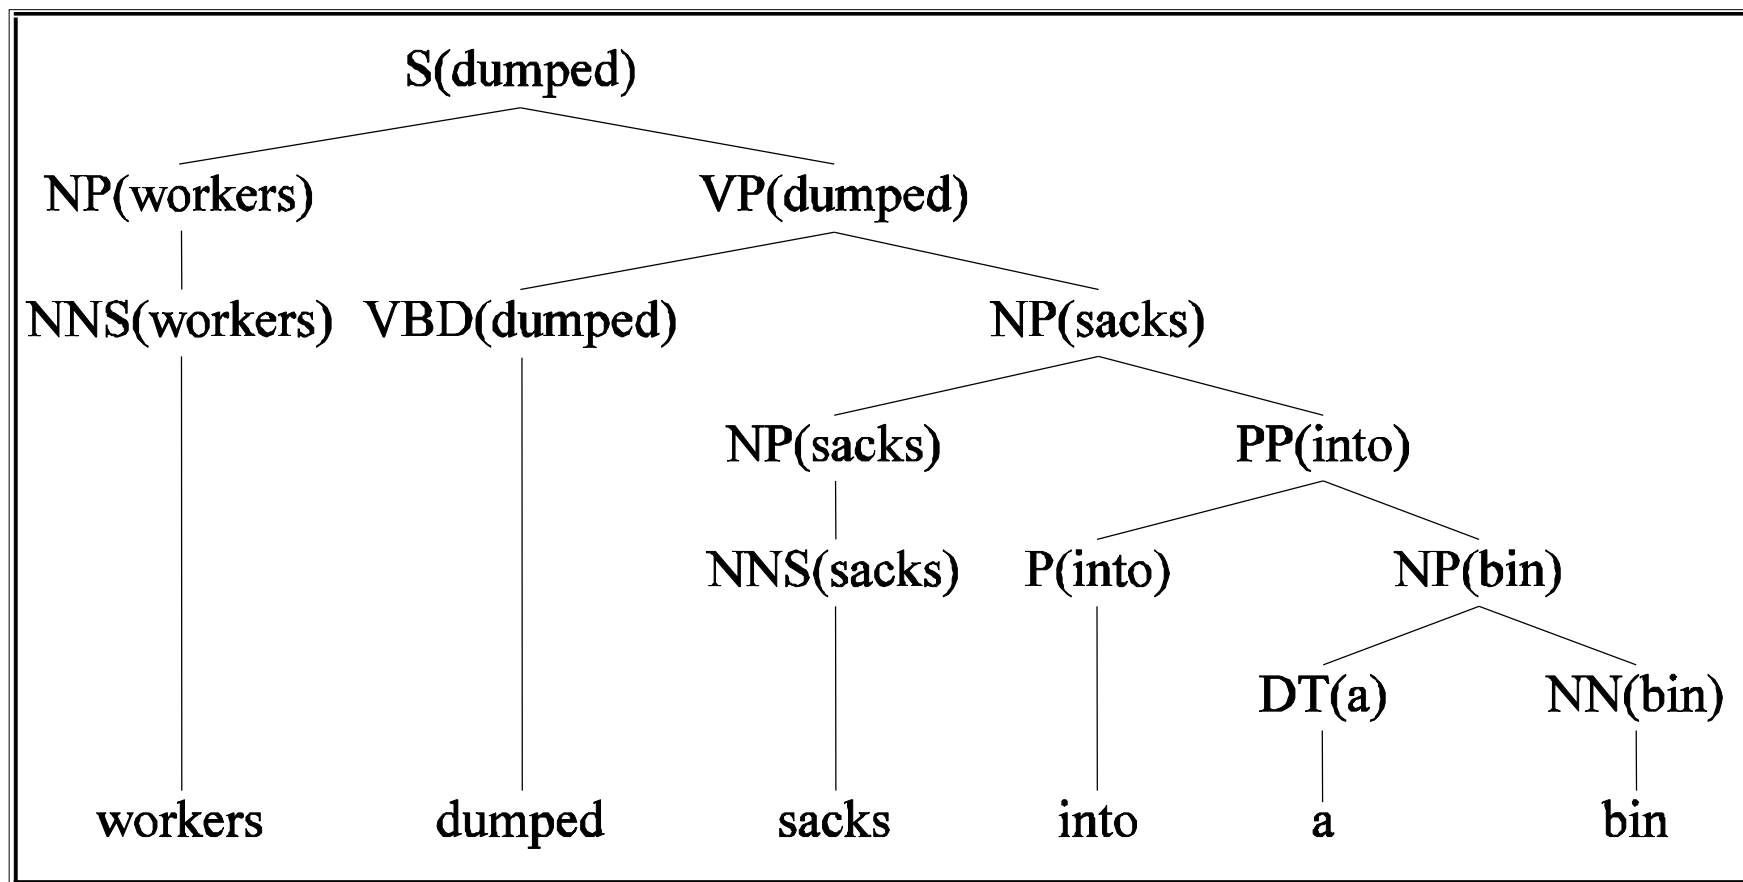
- To do that we' re going to make use of the notion of the **head 中心词** of a phrase
 - The head of an NP is its noun
 - The head of a VP is its verb
 - The head of a PP is its preposition(It's really more complicated than that but this will do.)



Example (sensible parse)



Example (incorrect parse)



How?

- We used to have
 - $VP \rightarrow V NP PP$ $P(\text{rule}|VP)$
 - That's the count of this rule divided by the number of VPs in a treebank
- Now we have
 - $VP(\text{dumped}) \rightarrow V(\text{dumped}) NP(\text{sacks}) PP(\text{in})$
 - $P(r|VP \wedge \text{dumped is the verb} \wedge \text{sacks is the head of the NP} \wedge \text{in is the head of the PP})$
 - Not likely to have significant counts in any treebank
 - Make independence assumptions to bread down each rule



Subcategorization

- Condition particular VP rules on their head...

SO

$r: VP \rightarrow V NP PP \quad P(r|VP)$

Becomes

$P(r \mid VP \wedge \text{dumped})$

What's the count?

How many times was this rule used with **dump**,
divided by the number of VPs that **dump** appears in
total

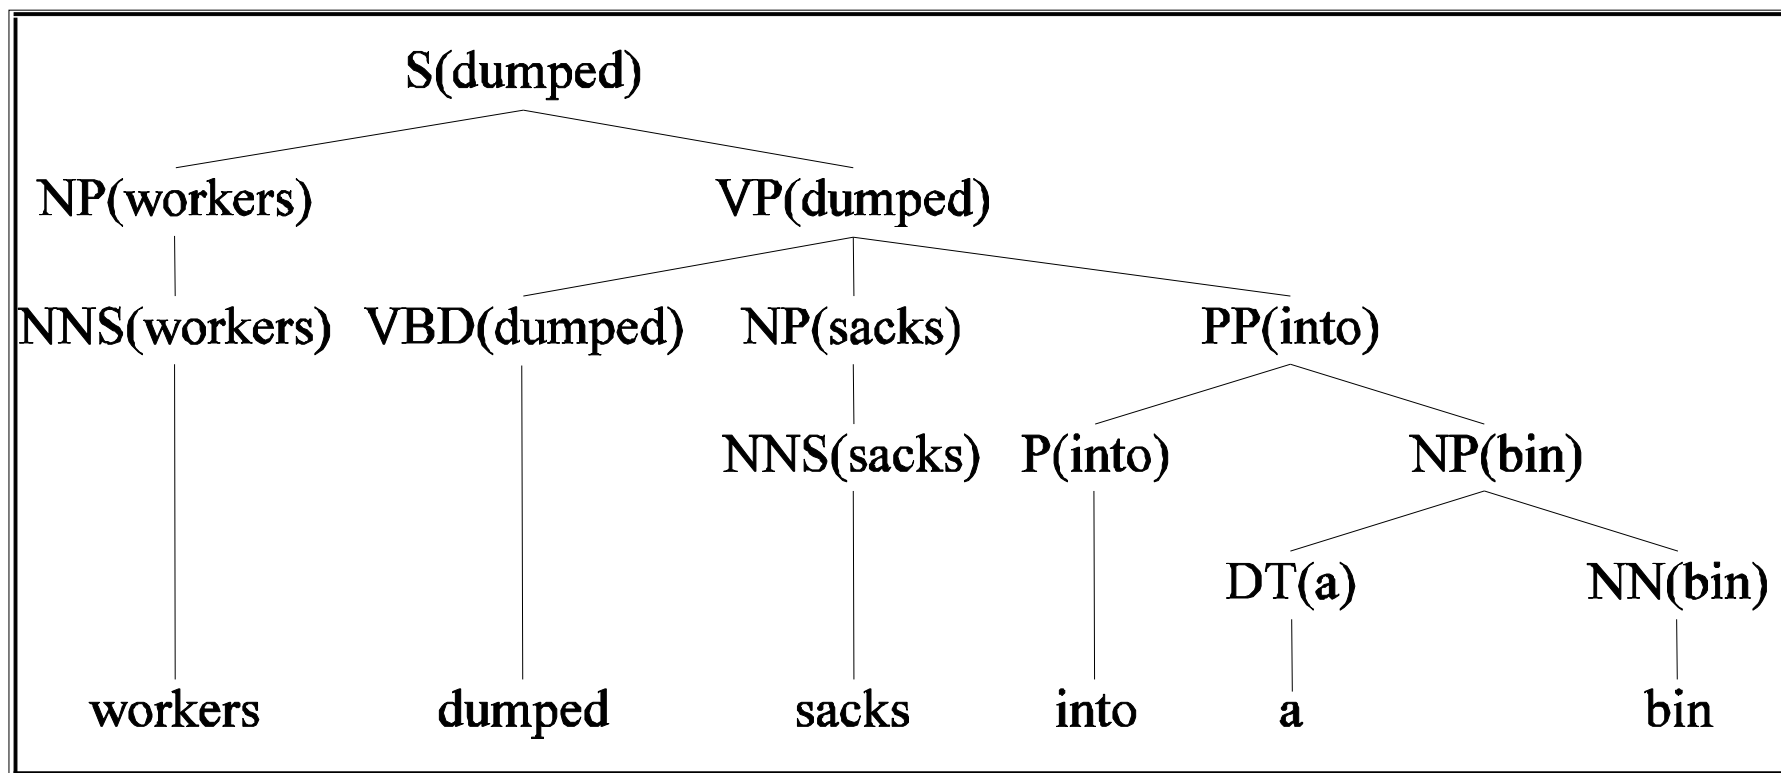


Preferences

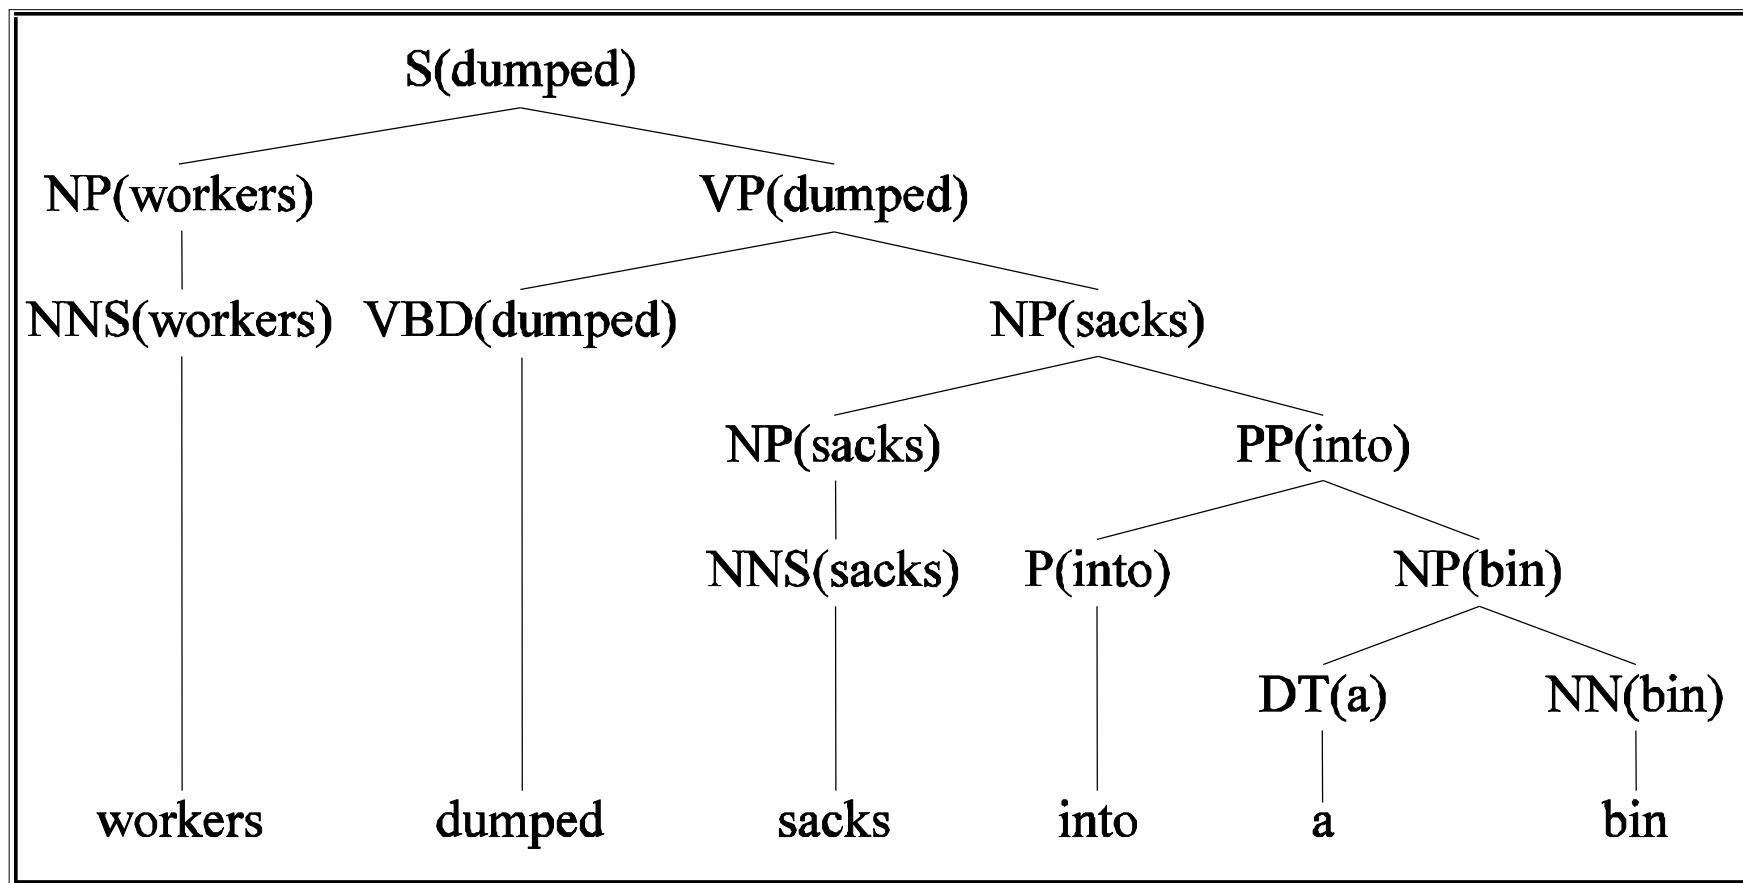
- Subcat captures the affinity between VP heads (verbs) and the VP rules they go with.
- What about the affinity between VP heads and the heads of the other daughters of the VP
- Back to our examples...



Example (right)



Example (wrong)



Preferences

- The issue here is the **attachment** of the PP. So the affinities we care about are the ones between **dumped** and **into** vs. **sacks** and **into**.
- So count the places where **dumped** is the head of a constituent that has a PP daughter with **into** as its head and normalize
- v.s. the situation where **sacks** is a constituent with **into** as the head of a PP daughter.





Preferences

- Consider the VPs
 - Ate spaghetti with gusto
 - Ate spaghetti with marinara
- The affinity of **gusto** for **eat** is much larger than its affinity for **spaghetti**
- On the other hand, the affinity of **marinara** for **spaghetti** is much higher than its affinity for **ate**



Preferences

- Note the relationship here is more distant and doesn't involve a headword since *gusto* and *marinara* aren't the heads of the PPs.



□ Probabilistic CFGs

□ Structural Dependencies

□ Parser Evaluation



Parser Evaluation

- See 14.7 (many parts are a repeat of 13.5.3 because chunking is partial parsing)
- Constituent-level evaluation
 - Sentence-level would be too coarse
- Cross brackets
 - number of brackets in the candidate parse which cross brackets in the treebank parse



Constituent Evaluation - Recall

$$\frac{\text{\# correct nodes in candidate parse}}{\text{\# nodes in treebank parse}}$$

Correct node = node in candidate parse
which:

- has same node label as in treebank
- spans the same words as in treebank



Constituent Eval - Precision


$$\frac{\text{\#correct nodes in candidate parse}}{\text{\# nodes in candidate parse}}$$

Correct node = node in candidate parse which:

- has same node label as in treebank
- spans the same words as in treebank



Example



Candidate	gold
X:a	X:a
Y:b	Z:b
Z:cd	V:cd
--	Y:bcd
W:abcd	W:abcd

Labeled Recall = 2/5; Labeled Precision = 2/4



Cross brackets

- Number of brackets in candidate parse that cross brackets in the treebank parse
 - e.g. treebank has ((X Y) Z) and candidate has (X (Y Z))
- Unlike precision/recall, this is an objective function to minimize



Drawbacks of PARSEVAL

- Rewards shallow/safe analyses better than those that make more claims but a few mistakes.
- Some "single" errors can hurt the score repeatedly, for example a single misplaced node may trigger multiple crossing brackets and incorrect nodes.
- Weights all nodes evenly, rather than making crucial semantical relations more important.



Metrics for Dependency Parsing

- Head Attachment Score (percent of nodes which are correctly attached to their parent)
- Label Precision (percent of nodes whose dependency labeled is predicted correctly)
- Labeled Attachment Score (percent of node for which both of the above are true)
- Branch Precision (percent of the Paths (from root to leaf) that are being classified correctly)
- Correct trees precision (percent of the sentences from the eval corpus which have been parsed flawlessly)

